

```
In [3]: # read the data into a pandas DataFrame
import pandas as pd
path = 'dataset/pima-indians-diabetes.data'
col_names = ['pregnant', 'glucose', 'bp', 'skin', 'insulin', 'bmi', 'pedigree', 'age', 'label']
pima = pd.read_csv(path, header=None, names=col_names)
```

```
In [5]: # print the first five data
pima.head()
```

Out[5]:

	pregnant	glucose	bp	skin	insulin	bmi	pedigree	age	label
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [6]: # define X and y
feature_cols = ['pregnant', 'insulin', 'bmi', 'age']
X = pima[feature_cols]
y = pima.label
```

```
In [8]: # split X and y into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

```
In [70]: import warnings
warnings.filterwarnings('ignore')
# train a logistic regression model on the training set
from sklearn.linear_model import LogisticRegression
```

```
logreg = LogisticRegression()  
logreg.fit(X_train, y_train)
```

```
Out[70]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=  
True,  
                             intercept_scaling=1, l1_ratio=None, max_iter=100,  
                             multi_class='warn', n_jobs=None, penalty='l2',  
                             random_state=None, solver='warn', tol=0.0001, verbos  
e=0,  
                             warm_start=False)
```

```
In [11]: # make class predication for the testing set  
y_pred_class = logreg.predict(X_test)
```

Classification accuracy: percentage of correct predictions

```
In [12]: # calculation accuracy  
from sklearn import metrics  
print(metrics.accuracy_score(y_test, y_pred_class))
```

```
0.6927083333333334
```

Null accuracy: accuracy that could be achieved by always predicting the most frequent class

```
In [13]: # examine the class distribution of the testing set (using a Pandas Ser  
ies Method)  
y_test.value_counts()
```

```
Out[13]: 0    130  
         1     62  
         Name: label, dtype: int64
```

```
In [14]: # calculate the percentage of ones  
y_test.mean()
```

```
Out[14]: 0.3229166666666667
```

```
In [16]: y_pred_class.mean()
```

```
Out[16]: 0.140625
```

```
In [17]: # calculate the percentage of zeros  
1 - y_test.mean()
```

```
Out[17]: 0.6770833333333333
```

```
In [18]: # calculate null accuracy (for binary classification problems coded as  
0/1)  
max(y_test.mean(), 1 - y_test.mean())
```

```
Out[18]: 0.6770833333333333
```

```
In [19]: # calculate null accuracy (for multi-class classification problems)  
y_test.value_counts() / len(y_test)
```

```
Out[19]: 0    0.677083  
1    0.322917  
Name: label, dtype: float64
```

```
In [20]: # calculate null accuracy (for multi-class classification problems)  
y_test.value_counts().head(1) / len(y_test)
```

```
Out[20]: 0    0.677083  
Name: label, dtype: float64
```

Comparing the true and predicted response values

```
In [21]: # print the first 25 true and predicted response  
print('True:', y_test.values[0:25])  
print('Pred:', y_pred_class[0:25])
```

```
True: [1 0 0 1 0 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 0 0]  
Pred: [0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

Conclusion: . Classification accuracy is the easiest classification metric to understand . But, it does not tell you the underlying distribution of response values . And, it does not tell you what "types" of errors your classifier is making

## Confusion matrix

Table that describes the performance of a classification model

```
In [22]: # Important: first argument is true values, second argument is predicted values  
print(metrics.confusion_matrix(y_test, y_pred_class))
```

```
[[118  12]  
 [ 47  15]]
```

- . Every observation in the testing set is represented in exactly one box.
- . It's a 2x2 matrix because there are 2 response classes.
- . The format shown here is not universal.

Basic terminology:

- . True Positives (TP): we correctly predicted that they do have diabetes
- . True Negatives (TN): we correctly predicted that they don't have diabetes
- . False Positives (FP): we incorrectly predicted that they do have diabetes (a "Type I error")
- . False Negatives (FN): we incorrectly predicted that they don't have diabetes (a "Type II error")

```
In [23]: # print the first 25 true and predicated response
print('True:', y_test.values[0:25])
print('Pred:', y_pred_class[0:25])
```

```
True: [1 0 0 1 0 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 0 0]
Pred: [0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

```
In [24]: # save confusion matrix and slice into four pieces
confusion = metrics.confusion_matrix(y_test, y_pred_class)
TP = confusion[1, 1]
TN = confusion[0, 0]
FP = confusion[0, 1]
FN = confusion[1, 0]
```

Metrics computed from a confusion matrix

Classification Accuracy: Overall, how often is the classifier correct?

```
In [25]: print((TP + TN) / float(TP + TN + FP + FN))
print(metrics.accuracy_score(y_test, y_pred_class))
```

```
0.6927083333333334
0.6927083333333334
```

Classification Error: Overall, how often is the classifier incorrect?

Also known as "Misclassification Rate"

```
In [26]: print((FP + FN) / float(TP + TN + FP + FN))
print(metrics.accuracy_score(y_test, y_pred_class))
```

```
0.3072916666666667
0.6927083333333334
```

Sensitivity: When the actual value is positive, how often is the prediction correct?

. How "sensitive" is the classifier to detecting positive instances?

. Also known as "True Positive Rate" or "Recall"

```
In [27]: print(TP / float(TP + FN))  
print(metrics.recall_score(y_test, y_pred_class))
```

0.24193548387096775

0.24193548387096775

Specificity: When the actual value is negative, how often is the prediction correct?

.How "specific" (or "selective") is the classifier in predicting positive instances?

```
In [31]: print(TN / float(TN + FP))
```

0.9076923076923077

False Positive Rate: When the actual value is negative, how often is the prediction incorrect?

```
In [32]: print(FP / float(TN + FP))
```

0.09230769230769231

Precision: When a positive value is predicted, how often is the prediction correct?

. How "precise" is the classifier when predicting positive instances?

```
In [34]: print(TP / float(TP + FP))  
print(metrics.precision_score(y_test, y_pred_class))
```

```
0.5555555555555556
0.5555555555555556
```

```
In [ ]: print(TP / float(TP + FP))
        print(metrics.precision_score(y_test, y_pred_class))
```

Many other metrics can be computed: F1 score, Matthews correlation coefficient, etc.

Conclusion:

- . Confusion matrix gives you a more complete picture of how your classifier is performing.
- . Also allows you to compute various classification metrics, and these metrics can guide your model selection.

Which metrics should you focus on?

- . Choice of metric depends on your business objective.
- . Spam filter (positive class is "spam"): Optimize for precision or specificity because false negatives (spam goes to the inbox) are more acceptable than false positives (non-spam is caught by the spam filter).
- . Fraudulent transaction detector (positive class is "fraud"): Optimize for sensitivity because false positives (normal transactions that are flagged as possible fraud) are more acceptable than false negatives (fraudulent transactions that are not detected).

Adjusting the classification threshold

```
In [35]: # print the first 10 predicted responses
        logreg.predict(X_test)[0:10]
```

```
Out[35]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 1], dtype=int64)
```

```
In [38]: # print the first 10 predicated probabilités of class membership  
logreg.predict_proba(X_test)[0:10, :]
```

```
Out[38]: array([[0.63247571, 0.36752429],  
                [0.71643656, 0.28356344],  
                [0.71104114, 0.28895886],  
                [0.5858938 , 0.4141062 ],  
                [0.84103973, 0.15896027],  
                [0.82934844, 0.17065156],  
                [0.50110974, 0.49889026],  
                [0.48658459, 0.51341541],  
                [0.72321388, 0.27678612],  
                [0.32810562, 0.67189438]])
```

```
In [39]: # print the first 10 predicated probabilités for class 1  
logreg.predict_proba(X_test)[0:10, 1]
```

```
Out[39]: array([0.36752429, 0.28356344, 0.28895886, 0.4141062 , 0.15896027,  
                0.17065156, 0.49889026, 0.51341541, 0.27678612, 0.67189438])
```

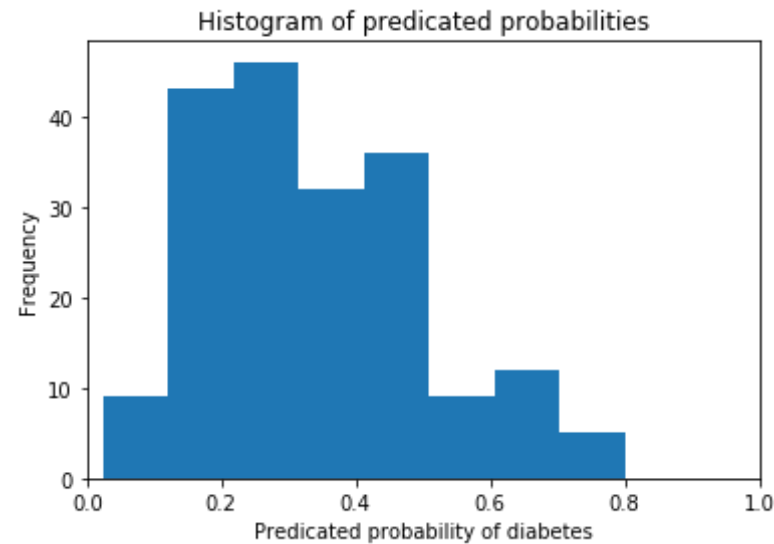
```
In [42]: y_pred_prob = logreg.predict_proba(X_test)[: , 1]
```

```
In [43]: # allow plots to appear in the notebook  
%matplotlib inline  
import matplotlib.pyplot as plt
```

```
In [44]: # histogram of predicated probabilités  
plt.hist(y_pred_prob, bins=8)  
plt.xlim(0, 1)  
plt.title('Histogram of predicated probabilities')  
plt.xlabel('Predicated probability of diabetes')  
plt.ylabel('Frequency')
```

```
Out[44]: Text(0, 0.5, 'Frequency')
```





Decrease the threshold for predicting diabetes in order to increase the sensitivity of the classifier

```
In [45]: # predicate diabetes if the predicted probability is greater than 0.3
from sklearn.preprocessing import binarize
y_pred_class = binarize([y_pred_prob], 0.3)[0]
```

```
In [46]: # print the first 10 predicted probabilities
y_pred_prob[0:10]
```

```
Out[46]: array([0.36752429, 0.28356344, 0.28895886, 0.4141062 , 0.15896027,
                0.17065156, 0.49889026, 0.51341541, 0.27678612, 0.67189438])
```

```
In [47]: # print the first 10 predicted class with the lower threshold
y_pred_class[0:10]
```

```
Out[47]: array([1., 0., 0., 1., 0., 0., 1., 1., 0., 1.])
```

```
In [49]: # previous confusion matrix (default threshold of 0.5)
print(confusion)
```

```
[[118 12]
 [ 47 15]]
```

```
In [50]: # new confusion matrix (threshold of 0.3)
print(metrics.confusion_matrix(y_test, y_pred_class))
```

```
[[80 50]
 [16 46]]
```

```
In [51]: # sensitivity has increased (used to be 0.24)
print(46 / float(46 + 16))
```

```
0.7419354838709677
```

```
In [52]: # specificity has decreased (used to be 0.91)
print(80 / float(80 + 50))
```

```
0.6153846153846154
```

Conclusion: . Threshold of 0.5 is used by default (for binary problems) to convert predicted probabilities into class predictions. . Threshold can be adjusted to increase sensitivity or specificity. . Sensitivity and specificity have an inverse relationship.


ROC Curves and Area Under the Curve (AUC)

Question: Wouldn't it be nice if we could see how sensitivity and specificity are affected by various thresholds, without actually changing the threshold?

Answer: Plot the ROC curve!

```
In [53]: # Important : first argument is true values, second argument is predicted probabilities
fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_prob)
plt.plot(fpr, tpr)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
```

```
plt.title('ROC curve for diabetes classifier')
plt.xlabel('False Positive Reate (1 - Specificity)')
plt.grid(True)
```



. ROC curve can help you to choose a threshold that balances sensitivity and specificity in a way that makes sense for your particular context. . You can't actually see the thresholds used to generate the curve on the ROC curve itself.

```
In [63]: # define a function that accepts a threshold and prints sensitivity and
         # specificity
         def evaluate_threshold(threshold):
             print('Sensitivity:', tpr[thresholds > threshold][-1])
             print('Specificity:', 1 - fpr[thresholds > threshold][-1])
```

```
In [64]: evaluate_threshold(0.5)

Sensitivity: 0.24193548387096775
Specificity: 0.9076923076923077
```

```
In [65]: evaluate_threshold(0.3)
```

Sensitivity: 0.7258064516129032  
Specificity: 0.6153846153846154

AUC is the percentage of the ROC plot that is underneath the curve:

```
In [67]: # Important: first argument is true values, second argument is predicted probabilities
print(metrics.roc_auc_score(y_test, y_pred_prob))

0.7245657568238213
```

. AUC is useful as a single number summary of classifier performance. . If you randomly chose one positive and one negative observation, AUC represents the likelihood that your classifier will assign a higher predicted probability to the positive observation. . AUC is useful even when there is high class imbalance (unlike classification accuracy).

```
In [69]: import warnings
warnings.filterwarnings('ignore')
# calculate cross-validated AUC
from sklearn.model_selection import cross_val_score
cross_val_score(logreg, X, y, cv=10, scoring='roc_auc').mean()
```

Out[69]: 0.7378233618233618

Confusion matrix advantages: . Allows you to calculate a variety of metrics . Useful for multi-class problems (more than two response classes)

ROC/AUC advantages: . Does not require you to set a classification threshold . Still useful when there is high class imbalance

In [ ]:

In [ ]: